



(a) headline cell, equal budget $N_{\text{train}} = 1.6 \times 10^5$



(b) the advantage holds at every budget



- Gaussian prior (sample covariance)
- GMM prior b^* (K by validation log-lik.)
- learned diffusion prior (proposed, V1)

- known-channel reference (R5)
- diffusion prior, best-val weights ($N = 3.2 \times 10^5$)